

Constraints

Data Integrity:

Maintaining and ensuring the accuracy and consistency of the data.

Ex:

- Students registered in a class are also in the Students relation.

(Foreign key)

- Each student has a different id (primary key)
- A student cannot register to more than 5 courses.
- etc. etc...

Integrity Constraints are a feature of DBMS to help guarantee integrity of the database.

Ex:

Primary keys
Attribute IS NOT NULL
Attribute IS UNIQUE
Foreign Key

} all types
of
constraints

Type of constraints:

- Attribute. Specified along declaration of attribute.
- tuple: Specified along table.
Applies to entire tuple

Table constraints checked everytime a tuple is inserted, updated or deleted.

- Database: Apply to entire DB.
Checked everytime the DB (any of its tuples) is inserted, updated, deleted.

Primary Key

- As attribute constraint (Key is only one attribute)

CREATE TABLE <name> (

⋮
attrname type PRIMARY KEY,

- As tuple constraint (multi attr. PK)

CREATE TABLE

⋮
declaration of attributes

⋮
PRIMARY KEY (List of attr). 2
⋮

UNIQUE

For other candidate keys you can use UNIQUE.

For one-attribute candidate keys:

⋮
attname type UNIQUE
⋮

Or more generally as tuple constraint:

⋮
UNIQUE (<att-list>)

↑ it can be one or more attributes.

UNIQUE is implied with Primary Key constraints.

NOT NULL

Only makes sense as an attribute constraint.

⋮
<attname> <type> NOT NULL;
⋮

Implicit for PKs, but UNIQUE att can be NULL

Referential Integrity: Foreign Key Constraint.

As attribute constraint.

⋮
attname type REFERENCES
 <relation>

As tuple constraint.

⋮
FOREIGN KEY (<attlist>) REFERENCES
 <relation> ↑
⋮ can be one or more attr.

Makes a FK constraint for attribute
attname to the primary key of <relation>

By default the reference is to the primary
key of the other table. But we can use
other attributes:

... REFERENCES <relation> (<attlist>)

But <attlist> must be declared UNIQUE

Assume

```
CREATE TABLE R(  
    a int PRIMARY KEY  
);  
CREATE TABLE S(  
    a int PRIMARY KEY,  
    FOREIGN KEY (a) REFERENCES R  
);
```

What if a tuple in R, referenced in S is deleted:

R	a
	3
	2
	5

S	a
	2
	5

What if we delete $\sigma_{a=5} R$?

What if we change in R $a=5$ to $a=6$?

4 options:

- 1) **CASCADE** Delete tuple in S too or update value in S to match new value in tuple of R
- 2) **RESTRICT** Deny if there are tuples that reference tuple being deleted. **Default!**
- 3) **SET NULL** Set the attribute(s) in the tuple that references to NULL and allow the delete or update of the tuple to proceed.
- 4) **SET DEFAULT** Replaces values of tuple in S with default values

Syntax

...
FOREIGN KEY (...) REFERENCES
...
ON { DELETE } { CASCADE
 UPDATE } { RESTRICT
 SET NULL
 SET DEFAULT }

Default

In insertions, attributes are set to NULL if not specified

Ex.

R(a,b,c)

INSERT INTO R(b) VALUES (5);
Rejected, the Primary Key (a) cannot be NULL.

INSERT INTO R(a) VALUES (3)
inserts:
(3, NULL, NULL) into R

We can change this behaviour:

⋮
<attname> <type> DEFAULT <value>

If not explicitly given, attribute is set to default value.

CHECK

Every time tuple is updated or tuple inserted a predicate is evaluated.

Operation fails unless predicate is true:

Ex:

year int CHECK (year > 1900),
gender char(1) CHECK
(gender IN ('F', 'M')),

Attribute CHECK

CHECK (a + b = 5)

↑

assuming both att.

tuple CHECK

are declared.

It can contain a Subquery (as any predicate in a selection):

customerid CHAR(10),

creditlimit REAL,

CHECK (creditlimit <=

SELECT sum(orders.amount)

FROM orders

WHERE orders.custid =
customerid).

Note how we use the attribute of the tuple being operated upon in the subquery. \vdots

$(creditLimit \leq$
 $SELECT \sum(orders.amount)$
 $FROM \text{orders}$
 $WHERE \text{orders.custid} =$
 $\underline{\text{customer id}}).$

value of current tuple \uparrow

This is a good use of correlated subqueries.

(In general avoid them because they tend to have horrible performance)

Altering Constraints

Every constraint gets a name.

We can give explicit names:

```
CONSTRAINT <name> <constraint>
```

Ex:

```
CONSTRAINT tablePK PRIMARY KEY (a)
```

Name becomes [↗] global!!

We can refer to it:

```
ALTER TABLE <tableName>
```

```
DROP CONSTRAINT <constraintName>
```

We can add constraints to an already created table:

```
ALTER TABLE ADD CONSTRAINT  
<name> <constraint>;
```

Ex:

```
ALTER TABLE R ADD CONSTRAINT  
myConst UNIQUE (a,b);
```

Assertions

An assertion is a predicate that must always be TRUE

```
CREATE ASSERTION <assertionName>  
CHECK ( <condition> )
```

Ex: No movie can have more than 100 roles.

```
CREATE ASSERTION NoMovieMore100roles  
CHECK (NOT EXISTS  
      (SELECT id FROM  
        Productions NATURAL  
        NATURAL JOIN Roles  
        WHERE att IS NULL  
        GROUP BY id  
        HAVING count(*) > 100));
```

At least one student passes each course:

```
CREATE ASSERTION AtLeastOne CHECK  
( 1 <= ALL  
  (SELECT count(*) FROM  
    Enrolled WHERE grade >= 50  
    GROUP BY cid)
```

True if $1 \leq$ every value in set.
true if set is empty.

Warning :

Assertions can slow a DB.

⇒ They are executed everytime the underlying relations change.

Postgres does not support them!!

Triggers

Triggers are events that are executed in response to an action.

- 1) An event "triggers" a trigger
- 2) Once started, a trigger checks a condition (optional). If not true trigger stops.
- 3) An action is executed

A trigger can be used to replace an event (INSTEAD OF)

Check can be done with the state of the database before or after action.

Trigger can be executed

- per each affected tuple, or
- one for all affected tuples.

A trigger can be used to

- Alter the behavior of an operation
- Do extra work.
- To abort the operation
- To update views (INSTEAD OF)

Syntax:

Postgres, Simplified

CREATE TRIGGER <name>

{	BEFORE	{	INSERT	}
	AFTER		DELETE	
	INSTEAD OF		UPDATE	
			UPDATE OF attname,...	

FOR EACH	ROW	}
	STATEMENT	

WHEN (condition)

EXECUTE PROCEDURE <name>(args)

Ex.

```
CREATE TRIGGER check-update  
  BEFORE UPDATE ON Accounts  
  FOR EACH ROW  
  WHEN (OLD.balance IS DISTINCT FROM  
        NEW.balance)  
  EXECUTE PROCEDURE check-account-update;
```

```
CREATE TRIGGER log-update  
  AFTER UPDATE  
  FOR EACH ROW  
  WHEN (OLD.* IS DISTINCT FROM NEW.*)  
  EXECUTE PROCEDURE log-acct-update;
```